

Harpua2001 Daily Routine -- Revamp Report

5-agent review of the daily automation pipeline, every fix applied and tested tonight, and the full-pipeline smoke-test result. Compiled 2026-08-03.

At a glance

sold_reconciler, sell_inventory_reprice, shipping_audit, listing_performance
4 fixes shipped

audited for shipping cost
1025 listings

found, 2 fixed live
3 mismatches

full pipeline smoke test
25/26 clean

Five specialist agents reviewed the daily pipeline end to end. Four concrete fixes shipped tonight, each tested against live data rather than just read and trusted. One real financial gap was caught (3 more shipping-undercharged listings, 2 already corrected) and one stale, silently-failing step was found and removed.

1. The 5-agent review -- what each lens found

Five independent specialists (ops/sequencing, card-selling strategy, marketing/customer-ops, promotions/pricing, risk/data-integrity) reviewed the same pipeline blind to each other's conclusions. Three of five independently flagged the same gap -- strong signal it was real.

- **Consensus finding (3/5 agents):** CollX has no inbound API, so a card sold on eBay never tells CollX it sold -- **sold_reconciler_agent.py** existed but never ran automatically. This is upstream of the oversell-guard cleanup work, not a duplicate of it.
- **Biggest structural gap (pricing lens):** repricing_agent.py can't touch listings posted via the Sell Inventory API (CollX imports) -- only sell_inventory_reprice.py can. It existed but was never previewed in the daily dry-run, so the digest's repricing section was blind to part of the catalog every single day.
- **Zero-detection risk (risk lens):** the 79-listing shipping-undercharge incident from 2026-08-02 had, and until tonight still had, no automated daily check of its own class of bug.
- **Fixed vs. flagged-not-ready:** pnl_agent.py and listing_performance_agent.py were both documented as broken and skipped daily -- worth fixing since they're the only real margin/search-rank visibility in the system. Separately, lot_generator_agent.py (the flagship AI Lot Generator) was found to cap lots at 25 cards in code, violating the standing 5-card rule -- flagged, not touched tonight.
- **Not yet actioned** (lower priority, no code changed): specifics_agent.py and feedback_agent.py additions, message_responder_agent.py / tracking_responder_agent.py for buyer-message visibility, a weekly email_campaign_agent.py dry-run, promotions_agent.py addition, and resolving a possible price_drops_agent vs. repricing_agent conflict on the same SKU. These are queued for a future pass.

2. Fixes shipped tonight, with real test results

sold_reconciler_agent.py -- added, pre-authorized

Added to Step 0a, right after the oversell guard. Confirmed safe to auto-apply by reading the code and running it live: **--apply only writes to the local state/linkage.db sqlite file** -- there is no eBay/marketplace write in either dry-run or --apply mode; both make the identical live GetMyeBaySelling read. Real run: 80 sold listings in a 30-day window, 1 matched to CollX linkage, 79 eBay-only/unlinked, no errors.

sell_inventory_reprice.py -- added to Step 1 dry-run

Added right after repricing_agent.py, since it depends on that run's output. A same-day test refined the original claim: repricing_agent's fresh dry-run produced 72 apply-targets, but only **9 (12.5%)** were actual CollX/Inventory-API SKUs needing this fallback -- the other 63 were ordinary Trading listings repricing_agent could already reach directly. The 97%-of-catalog figure describes overall catalog composition (how many listings were posted via the newer API path), not what share of any single day's repricing targets needs the fallback -- both numbers are now documented so the daily summary doesn't overstate the gap.

listing_performance_agent.py -- fixed

The skill doc blamed a schema-drift crash. On inspection, that guard had already been patched separately at some point -- the real, still-live bug was a Sell Analytics API date-range error: requesting a window ending "today" gets rejected as "end date in the future" because Analytics data lags live eBay by about a day. Fixed by ending the window at yesterday. Verified with a real run: 85,661 impressions / 718 clicks / 200 listings analyzed, no errors.

pnl_agent.py -- confirmed working, no fix needed

Tested directly rather than trusting the skill doc's "broken" label -- it ran clean on the first try, correctly surfacing negative-margin listings (several Icon Collection Mahomes cards losing \$10-12 each after fees). Whatever schema-drift issue existed previously no longer reproduces.

shipping_audit_agent.py -- new agent, built and run live

New report-only SRE-style gate: live Trading API GetItem per active listing (the snapshot doesn't carry shipping fields), cached 7 days so only new/stale listings get re-checked daily. First cold run against all 1,025 active listings took 18.9 minutes with an empty cache; subsequent daily runs will be fast. **Found 3 real mismatches** -- the original 79-listing fix from 2026-08-02 wasn't fully complete:

- Frank Thomas Heavy Lumber Bat Relic Auto (\$152.50) -- USPSPriority @ \$0.00 (should be US_eBayStandardEnvelope @ \$1.32). **Not fixed** -- Reviseltem rejected it for an unclear reason; ruled out Best Offer, an active bid, and ending-within-12-hours via follow-up GetItem/GetBestOffers checks. Needs a manual fix in Seller Hub.
- Maxx Crosby Signature Class Round 4 Pick 4 (\$14.99) -- was USPSFirstClass @ \$0.99. **Fixed** via Reviseltem, confirmed live at \$1.32/standard envelope.
- Javonte Williams Select Concourse (\$4.49) -- was USPSFirstClass @ \$0.99. **Fixed** via Reviseltem, confirmed live.

Bonus catch: a stale, silently-failing pipeline step

The full pipeline smoke test (below) surfaced one failure: vault_eligibility.py. Root cause wasn't a regression from tonight's changes -- the script was **deleted from the repo on 2026-05-28** ("Nav cleanup: drop Vault + Whatnot") but never removed from the daily skill's agent list. This step has been failing every single day for over two months without anyone noticing, since a single agent failure doesn't stop the rest of the pipeline. Removed from the skill.

3. Full-pipeline smoke test

Ran all 26 Step 1 agents end to end, in the exact revised order, to catch any breakage from reordering rather than testing each agent in isolation. **25 of 26 ran clean** on the first pass, including the two new sequencing dependencies (shipping_audit_agent.py running early with a warm cache, sell_inventory_reprice.py correctly consuming repricing_agent.py's just-written plan). The one failure (vault_eligibility.py, above) was diagnosed, fixed by removing the stale reference, and the pipeline is now clean end to end.

4. Current state / what's still open

- **Needs your action:** the Frank Thomas \$152.50 bat relic auto still has wrong shipping (USPSPriority @ \$0.00) -- fix it directly in Seller Hub, since the automated Reviseltem path rejected it for a reason that wasn't Best Offer, a bid, or an ending-soon lock.
- **Daily pipeline is ready to run** with tonight's four additions/fixes plus the vault_eligibility.py removal -- no further setup needed.
- **Queued for later, not done tonight:** specifics_agent.py / feedback_agent.py additions, message_responder_agent.py / tracking_responder_agent.py for buyer-message visibility, a weekly email_campaign_agent.py preview, promotions_agent.py addition, resolving the price_drops_agent/repricing_agent overlap, and fixing the 5-card cap bug in lot_generator_agent.py before it's ever added to the daily routine.

Everything above was tested against live data, not just read and assumed -- every fix has a real run's output behind it, and the smoke test caught a two-month-old silent failure the original skill doc never mentioned.